Notes on the accuracy of naive summation

S. G. Johnson, MIT Course 18.335

February 8, 2019

1 Naive summation

In these notes, we analyze the floating-point error involved in summing n numbers, i.e. in computing the function $f(x) = \sum_{i=1}^{n} x_i$ for $x \in \mathbb{F}^n$ ($\mathbb{F}$ being the set of floating-point numbers), where the sum is done in the most obvious way, in sequence. In pseudocode:

```
sum = 0
for i = 1 to n
    sum = sum + x_i
f(x) = sum
```

A much more complete analysis of summation can be found in Higham (1993) [1]. Perhaps confusingly, this naive algorithm is called "recursive" summation, in reference to the inductive version of the definition below, although most computer programs would implement this with a loop (with the exception of Lisp programmers using tail recursion).

For analysis, it is a bit more convenient to define the process inductively:

$$\begin{aligned} s_0 &= 0 \\ s_k &= s_{k-1} + x_k \text{ for } 0 < k \leq n, \end{aligned}$$

with $f(x) = s_n$. (The intermediate values s_k are known as "partial" sums.) When we implement this in floating-point arithmetic, we get the function $\tilde{f}(x) = \tilde{s}_n$, where $\tilde{s}_k = \tilde{s}_{k-1} \oplus x_k$, with $\oplus$ denoting (correctly rounded) floating-point addition.

2 An upper bound on the error

We can easily prove an upper bound on the errors accumulated by the floating-point implementation of this algorithm:

$$|\tilde{f}(x) - f(x)| \leq n\epsilon_{\text{machine}} \sum_{i=1}^{n} |x_i| + O(\epsilon_{\text{machine}}^2) = n\epsilon_{\text{machine}}|f(x)| \left[\frac{\sum_{i=1}^{n} |x_i|}{|\sum_{i=1}^{n} x_i|} \right].$$

This means that the *relative error* in the sum is bounded above by

$$\frac{|\tilde{f}(x) - f(x)|}{|f(x)|} \leq n\epsilon_{\text{machine}} \left[\frac{\sum_{i=1}^{n} |x_i|}{|\sum_{i=1}^{n} x_i|} \right].$$

The $[\cdots]$ factor is what we will eventually call the **condition number** of the summation problem, a term that we we will define precisely later in 18.335. In the special case of summing *nonnegative* values $x_i \geq 0$, the $[\cdots]$ term is $= 1$, and we find that the relative error grows **at worse linearly** with the problem size n.

To prove this, we first prove the lemma:

$$\tilde{f}(x) = \sum_{i=1}^{n} x_i \prod_{k=i}^{n} (1 + \epsilon_k),$$

where $\epsilon_1 = 0$ and the other ϵ_k satisfy $|\epsilon_k| \leq \epsilon_{\text{machine}}$, by induction on n.

- For $n = 1$, it is trivial with $\epsilon_1 = 0$.

- Now for the inductive step. Suppose $\tilde{s}_{n-1} = \sum_{i=1}^{n-1} x_i \prod_{k=i}^{n-1}(1 + \epsilon_k)$. Then $\tilde{s}_n = \tilde{s}_{n-1} \oplus x_n = (\tilde{s}_{n-1} + x_n)(1 + \epsilon_n)$ where $|\epsilon_n| < \epsilon_{\text{machine}}$ is guaranteed by floating-point addition. The result follows by inspection: the previous terms are all multiplied by $(1 + \epsilon_n)$, and we add a new term $x_n(1 + \epsilon_n)$.

Now, let us multiply out the terms:

$$(1 + \epsilon_i) \cdots (1 + \epsilon_n) = 1 + \sum_{k=i}^{n} \epsilon_k + (\text{products of } \epsilon) = 1 + \delta_i,$$

where the products of ϵ_k terms are $O(\epsilon_{\text{machine}}^2)$, and hence

$$|\delta_i| \leq \sum_{k=i}^{n} |\epsilon_k| + O(\epsilon_{\text{machine}}^2) \leq n\epsilon_{\text{machine}} + O(\epsilon_{\text{machine}}^2).$$

Now we have: $\tilde{f}(x) = f(x) + (x_1 + x_2)\delta_2 + \sum_{i=3}^{n} x_i \delta_i$, and hence (by the triangle inequality):

$$|\tilde{f}(x) - f(x)| \leq |x_1| |\delta_2| + \sum_{i=2}^{n} |x_i| |\delta_i|.$$

Hence $|\tilde{f}(x) - f(x)| \leq n\epsilon_{\text{machine}} \sum_{i=1}^{n} |x_i|$ from the $|\delta_i|$ bound above.

Note: This does *not* correspond to a proof of forwards stability (defined soon in 18.335), since we have only shown that $|\tilde{f}(x) - f(x)| = \|x\| O(\epsilon_{\text{machine}})$, which is different from $|\tilde{f}(x) - f(x)| = |f(x)| O(\epsilon_{\text{machine}})$ *unless* all the x_i are ≥ 0! Note that our $O(\epsilon_{\text{machine}})$ is uniformly convergent in x, however (that is, the coefficient of $\epsilon_{\text{machine}}$ is independent of x, although it depends on n).

3 Average errors

In fact, the analysis above is typically too pessimistic, because the individual errors ϵ_k are typically of *different signs*, and in particular can usually be though of as random numbers, because the last few digits of typical inputs x_i are usually random noise. For uniform random ϵ_k, since δ_i is the sum of $(n - i + 1)$ random variables with variance $\sim \epsilon_{\text{machine}}$, it follows from the usual properties of **random walks** that the mean $|\delta_i|$ has magnitude $\sim \sqrt{n - i + 1} O(\epsilon_{\text{machine}}) \leq \sqrt{n} O(\epsilon_{\text{machine}})$. Hence we typically expect

$$\text{root mean square } |\tilde{f}(x) - f(x)| = O\left(\sqrt{n} \epsilon_{\text{machine}} \sum_{i=1}^{n} |x_i|\right),$$

i.e. rms errors that grow $\sim \sqrt{n}$.

This sounds good, but in fact there are summation algorithms that do **much better**. The algorithm for Julia's built-in `sum` function, for example, is **pairwise summation**, which has $O(\log n)$ worst-case and $O(\sqrt{\log n})$ average-case errors [1], while having about the same performance as naive summation.

References

[1] Nicholas J. Higham, "The accuracy of floating point summation," *SIAM Journal on Scientific Computing* **14**, pp. 783–799 (1993).