


Experiments with Cache-Oblivious Matrix Multiplication for 18.335

Steven G. Johnson
MIT Applied Math

platform: 2.66GHz Intel Core 2 Duo,
GNU/Linux + gcc 4.1.2 (-O3) (64-bit), double precision

(optimal) Cache-Oblivious Matrix Multiply



divide and conquer:

divide C into 4 blocks
compute block multiply recursively

achieves optimal $\Theta(n^3/\sqrt{Z})$ cache complexity

A little C implementation (~25 lines)

```
/* C = C + AB, where A is m x n, B is n x p, and C is m x p, in
row-major order. Actually, the physical size of A, B, and C
are m x fda, n x fdb, and m x fdc, but only the first n/p/p
columns are used, respectively. */
void add_matmul_rec(const double *A, const double *B, double *C,
int m, int n, int p, int fda, int fdb, int fdc)
{
    if (m+n+p <= 48) { /* <= 16x16 matrices "on average" */
        int i, j, k;
        for (i = 0; i < m; ++i)
            for (k = 0; k < p; ++k) {
                double sum = 0;
                for (j = 0; j < n; ++j)
                    sum += A[i*fda + j] * B[j*fdb + k];
                C[i*fdc + k] += sum;
            }
    }
    else { /* divide and conquer */
        int m2 = m/2, n2 = n/2, p2 = p/2;
        add_matmul_rec(A, B, C, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A+m2, B+m2*fdb, C, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A, B+p2, C+p2, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A+m2, B+p2*fdb, C+p2, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A+m2*fda, B, C+m2*fdc, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A+m2*fda, B+p2, C+m2*fdc+p2, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A+m2*fda, B+p2*fdb, C+m2*fdc, m2, n2, p2, fda, fdb, fdc);
        add_matmul_rec(A+m2*fda, B+p2*fdb, C+m2*fdc+p2, m2, n2, p2, fda, fdb, fdc);
    }
}

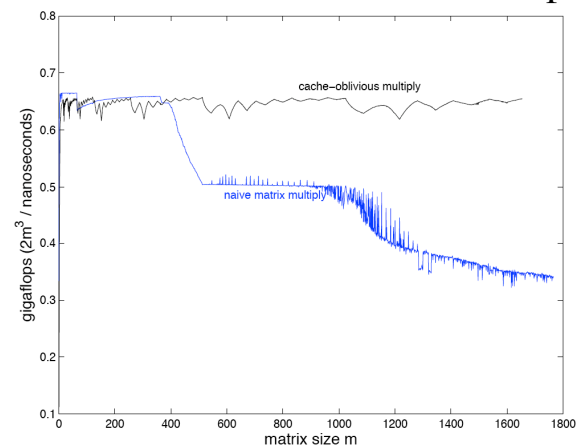
void matmul_rec(const double *A, const double *B, double *C,
int m, int n, int p)
{
    memset(C, 0, sizeof(double) * m*p);
    add_matmul_rec(A, B, C, m, n, p, p, p, p);
}
```

note: base case is ~16x16


*recursing down to 1x1
would kill performance
(1 function call per element,
no register re-use)*

dividing C into 4
— note that, instead, for
very non-square matrices,
we might want to divide
 C in 2 along longest axis

No Cache-based Performance Drops!



...but absolute performance still sucks



Registers .EQ. Cache

- The registers (~100) form a very small, almost ideal cache
 - Three nested loops is not the right way to use this “cache” for the same reason as with other caches
- Need long blocks of unrolled code: load blocks of matrix into local variables (= registers), do matrix multiply, write results
 - Loop-free blocks = many optimized hard-coded base cases of recursion for different-sized blocks ... often automatically generated (ATLAS)
 - Unrolled $n \times n$ multiply has $(n^3)!$ possible code orderings — compiler cannot find optimal schedule (NP hard) — cache-oblivious scheduling can help (c.f. FFTW), but ultimately requires some experimentation (automated in ATLAS)