

18.335 Problem Set 3

Due Friday, 15 March 2019.

Problem 1: QR and orthogonal bases

- (a) Trefethen, problem 10.4.
- (b) Prove that $A = QR$ and $B = RQ$ have the same eigenvalues, assuming A is a square matrix. Then do a little experiment: Construct a random 5×5 real-symmetric matrix in Julia via `X=rand(5,5); A = X' + X`. Use `QR = qr(A)` (do using `LinearAlgebra` first) to compute the QR factorization of A , and then compute $B = QR.R * QR.Q$. Then find the QR factorization $B = Q'R'$, and compute $R'Q'$...repeat this process until the matrix converges (maybe writing a loop and/or a function). From what it converges to, suggest a procedure to compute the eigenvalues and eigenvectors of a matrix (no need to prove that it converges in general—we will discuss this in class).
- (c) Trefethen, problem 28.2.

Problem 2: Schur fine

In class, we will show that any square $m \times m$ matrix A can be factorized as $A = QTQ^*$ (the *Schur factorization*), where Q is unitary and T is an upper-triangular matrix (with the same eigenvalues as A , since the two matrices are similar).

- (a) A is called “normal” if $AA^* = A^*A$. Show that this implies $TT^* = T^*T$. From this, show that T must be diagonal. Hence, any normal matrix (e.g. unitary or Hermitian matrices) must be unitarily diagonalizable. Hint: consider the diagonal entries of TT^* and T^*T , starting from the (1,1) entries and proceeding diagonally downwards by induction.
- (b) Given the Schur factorization of an arbitrary A (not necessarily normal), describe an algorithm to find the eigenvalues and eigenvectors of A , assuming for simplicity that all the eigenvalues are distinct. The flop count should be asymptotically $Km^3 + O(m^2)$; give the constant K .

Problem 3: Caches and backsubstitution

In this problem, you will consider the impact of caches (again in the ideal-cache model from class) on the problem of *backsubstitution*: solving $Rx = b$ for x , where R is an $m \times m$ upper-triangular matrix (such as might be obtained by Gaussian elimination). The simple algorithm you probably learned in previous linear-algebra classes (and reviewed in the book, lecture 17) is (processing the rows from bottom to top):

$$\begin{aligned} x_m &= b_m / r_{mm} \\ \text{for } j &= m-1 \text{ down to } 1 \\ x_j &= (b_j - \sum_{k=j+1}^m r_{jk}x_k) / r_{jj} \end{aligned}$$

Suppose that X and B are $m \times n$ matrices, and we want to solve $RX = B$ for X —this is equivalent to solving $Rx = b$ for n different right-hand sides b (the n columns of B). One way to solve the $RX = B$ for X is to apply the standard backsubstitution algorithm, above, to each of the n columns in sequence.

- (a) Give the asymptotic cache complexity $Q(m, n; Z)$ (in asymptotic Θ notation, ignoring constant factors) of this algorithm for solving $RX = B$.
- (b) Suppose $m = n$. Propose an algorithm for solving $RX = B$ that achieves a better asymptotic cache complexity (by cache-aware/blocking or cache-oblivious algorithms, your choice). Can you gain the factor of $1/\sqrt{Z}$ savings that we showed is possible for square-matrix multiplication?